

# Graphical Models

## Lecture 12: Belief Update Message Passing

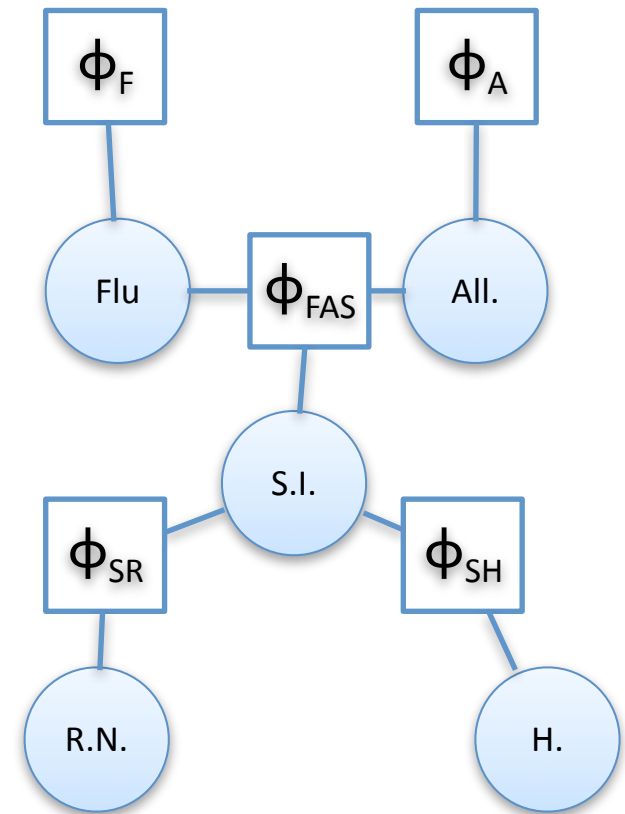
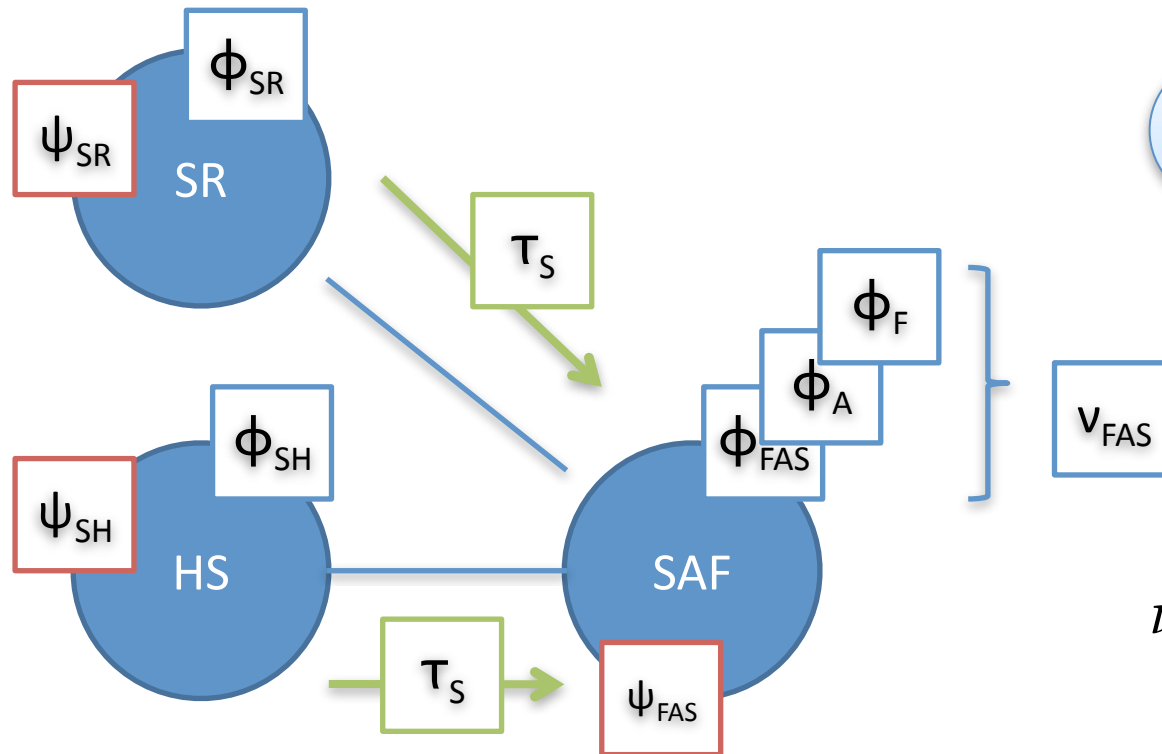
Andrew McCallum  
mccallum@cs.umass.edu

# Today's Plan

- Quick Review: Sum Product Message Passing (also known as “Shafer-Shenoy”)
- Today: Sum Product **Divide** Message Passing (also known as “belief update” message passing “Lauritzen-Spiegelhalter” and “belief propagation”)
- Mathematically equivalent, but different intuitions.
- Moving toward approximate inference.

# Quick Review

- $\{H, R, S, A, F\}$



$$\nu_j(\mathbf{C}_j) = \prod_{\phi \in \Phi_j} \phi$$

# Message Passing (One Root)

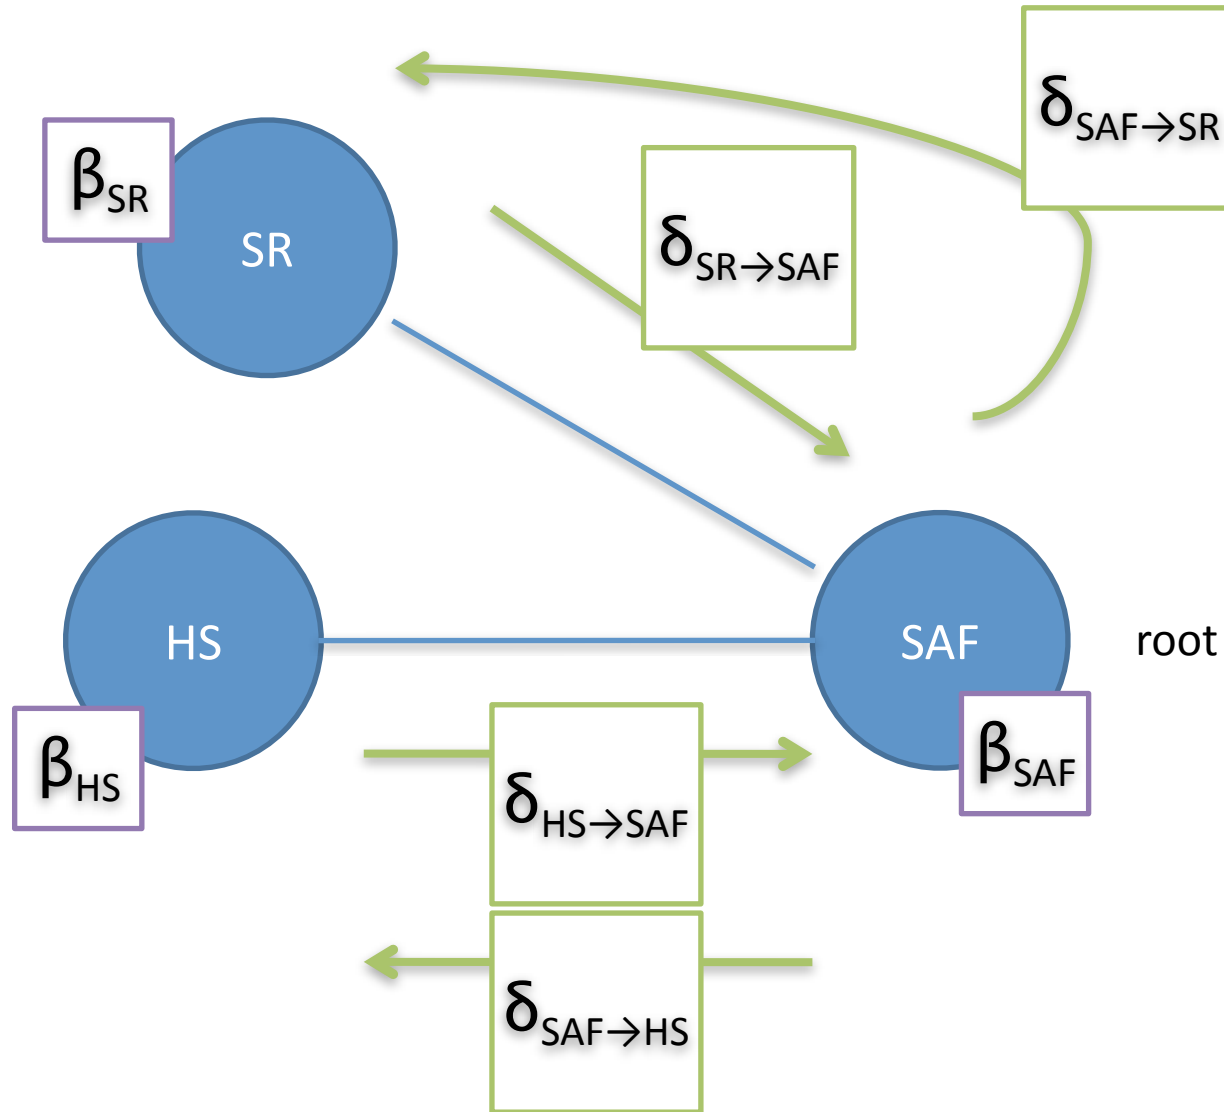
- Input: clique tree  $\mathcal{T}$ , factors  $\Phi$ , root  $\mathbf{C}_r$
- For each clique  $\mathbf{C}_i$ , calculate  $\nu_i$
- While  $\mathbf{C}_r$  is still waiting on incoming messages:

– Choose a  $\mathbf{C}_i$  that *has* received all of its incoming messages.

– Calculate and send the

$$\begin{aligned} \text{message from } \mathbf{C}_i \text{ to } \mathbf{C}_{\text{upstream-neighbor}(i)} &= \nu_i \prod_{k \in \text{Neighbors}_i \setminus \{j\}} \delta_{k \rightarrow i} \\ &= \sum_{\mathbf{X} \setminus \mathbf{C}_i} \prod_{\phi \in \Phi} \phi \\ &= Z \cdot P(\mathbf{C}_r) \end{aligned}$$

# Different Roots; Same Messages



# Sum-Product Message Passing

- Each clique tree vertex  $\mathbf{C}_i$  passes messages to each of its neighbors once it's ready to do so.
- At the end, for all  $\mathbf{C}_i$ :

$$\beta_i = \nu_i \prod_{k \in \text{Neighbors}_i} \delta_{k \rightarrow i}$$


- This is the unnormalized marginal for  $\mathbf{C}_i$ .

# Calibrated Clique Tree

- Two adjacent cliques  $\mathbf{C}_i$  and  $\mathbf{C}_j$  are calibrated

when:

$$\begin{aligned} \sum_{\mathbf{C}_i \setminus \mathbf{S}_{i,j}} \beta_i &= \sum_{\mathbf{C}_j \setminus \mathbf{S}_{i,j}} \beta_j \\ &= \mu_{i,j}(\mathbf{S}_{i,j}) \end{aligned}$$



# Calibrated Clique Tree as a

- Original (unnormalized) factor model and calibrated clique tree represent the same (unnormalized) measure:

$$\prod_{\phi \in \Phi} \phi = \frac{\prod_{C \in \text{Vertices}(\mathcal{T})} \beta_C}{\prod_{S \in \text{Edges}(\mathcal{T})} \mu_S}$$

Diagram illustrating the relationship between the original factor model and the calibrated clique tree representation:

- The numerator  $\prod_{\phi \in \Phi} \phi$  represents the **unnormalized Gibbs distribution from original factors  $\Phi$** .
- The denominator  $\prod_{S \in \text{Edges}(\mathcal{T})} \mu_S$  represents the **sepset beliefs**.
- The product  $\prod_{C \in \text{Vertices}(\mathcal{T})} \beta_C$  represents the **clique beliefs**.



# Inventory of Factors

- original factors  $\phi$
- initial potentials  $v$
- messages  $\delta$
- intermediate factors  $\psi$  (no longer explicit)
- clique beliefs  $\beta$
- sepset beliefs  $\mu$

# Inventory of Factors

- original factors  $\phi$
- initial potentials  $v$
- messages  $\delta$
- intermediate factors  $\psi$  (no longer explicit)
- clique beliefs  $\beta$
- sepset beliefs  $\mu$

New algorithm  
collapses  
everything into  
**beliefs!**

# Another Operation: Factor Division

- 0 / 0 is defined to be 0
- a / 0 is undefined when a > 0

A	B	C	$\phi_1(A, B, C)$
0	0	0	7000
0	0	1	30
0	1	0	5
0	1	1	500
1	0	0	100
1	0	1	1
1	1	0	10
1	1	1	1000

/

B	C	$\phi_2(B, C)$
0	0	100
0	1	1
1	0	2
1	1	100

=

A	B	C	$\phi_3(A, B, C)$
0	0	0	70
0	0	1	30
0	1	0	2.5
0	1	1	5
1	0	0	1
1	0	1	1
1	1	0	5
1	1	1	10

# Messages

- When computing the message  $\delta_{i \rightarrow j}$  from  $i$  to  $j$ , we multiply together all incoming messages to  $i$  *except* the one from  $j$  to  $i$ ,  $\delta_{j \rightarrow i}$ .
- Alternative: multiply *all* messages, and *divide* out the one from  $j$  to  $i$ .

$$\beta_i = \nu_i \prod_{k \in \text{Neighbors}_i} \delta_{k \rightarrow i}$$

$$\delta_{i \rightarrow j} = \sum_{C_i \setminus S_{i,j}} \nu_i \prod_{k \in \text{Neighbors}_i \setminus \{j\}} \delta_{k \rightarrow i}$$

$$\delta_{i \rightarrow j} = \frac{\sum_{C_i \setminus S_{i,j}} \beta_i}{\delta_{j \rightarrow i}}$$

# Key Idea

- We can “forget” the initial potentials  $v$ .
- We do not need to calculate the messages  $\delta$  explicitly.
- Store a partially calculated  $\beta$  on each vertex and a partially calculated  $\mu$  on each edge; *update* whenever new information comes in.

# A Single Belief Update

- At any point in the algorithm:

$$\begin{aligned}\sigma_{i \rightarrow j} &\leftarrow \sum_{\mathcal{C}_i \setminus \mathcal{S}_{i,j}} \beta_i \\ \beta_j &\leftarrow \beta_j \times \frac{\sigma_{i \rightarrow j}}{\mu_{i,j}} \\ \mu_{i,j} &\leftarrow \sigma_{i \rightarrow j}\end{aligned}$$

# Belief Update Message Passing

- Maintain beliefs at each vertex ( $\beta$ ) and edge ( $\mu$ ).
- Initialize each  $\beta_i$  to  $v_i$ .
- Initialize each  $\mu_{i,j}$  to **1**.
- Pass belief update messages.

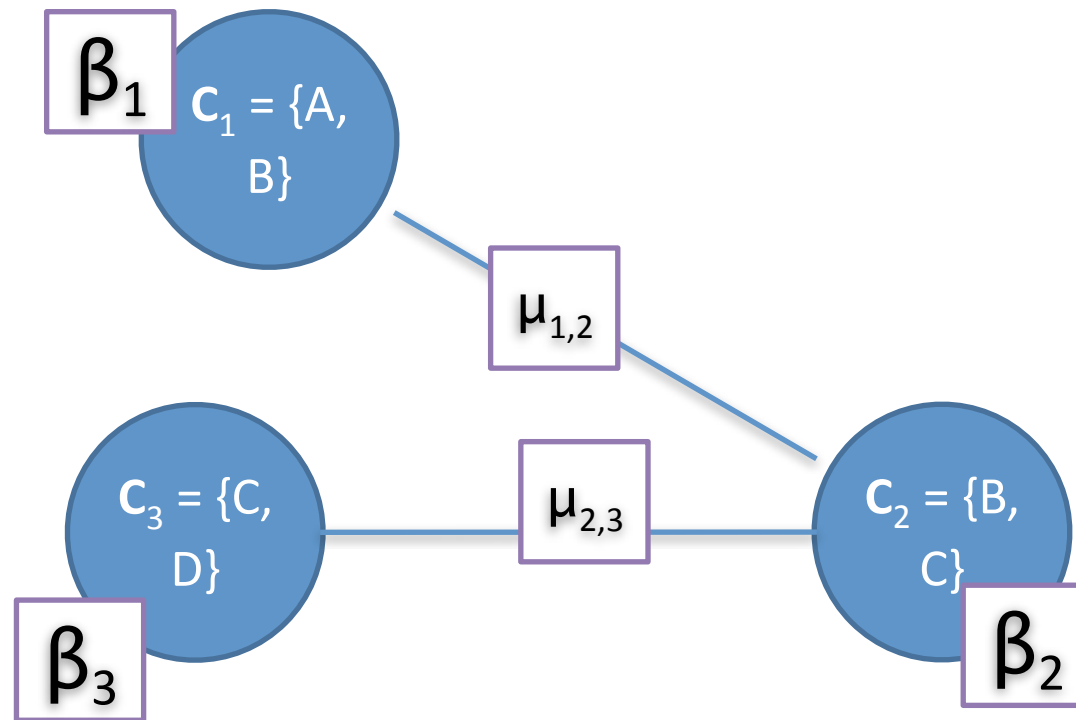
$$\sigma_{i \rightarrow j} \leftarrow \sum_{C_i \setminus S_{i,j}} \beta_i$$

$$\beta_j \leftarrow \beta_j \times \frac{\sigma_{i \rightarrow j}}{\mu_{i,j}}$$

$$\mu_{i,j} \leftarrow \sigma_{i \rightarrow j}$$

# Three Clique Example

$$\begin{aligned}\sigma_{i \rightarrow j} &\leftarrow \sum_{C_i \setminus S_{i,j}} \beta_i \\ \beta_j &\leftarrow \beta_j \times \frac{\sigma_{i \rightarrow j}}{\mu_{i,j}} \\ \mu_{i,j} &\leftarrow \sigma_{i \rightarrow j}\end{aligned}$$





# Worries

- Does the order of the messages matter?
- What if we pass the same message twice?
- What if we pass a message based on partial information?

# Claims

- At convergence, we will have a calibrated clique tree.

$$\begin{aligned} \sum_{C_i \setminus S_{i,j}} \beta_i &= \sum_{C_j \setminus S_{i,j}} \beta_j \\ &= \mu_{i,j}(S_{i,j}) \end{aligned}$$

- Invariant: throughout the algorithm:

$$\prod_{\phi \in \Phi} \phi = \prod_{C \in \text{Vertices}(\mathcal{T})} \nu_C = \frac{\prod_{C \in \text{Vertices}(\mathcal{T})} \beta_C}{\prod_{S \in \text{Edges}(\mathcal{T})} \mu_S}$$

# Equivalence

- Sum product message passing and sum product divide message passing lead to the same result: a calibrated clique tree.
  - *SumProduct* lets you calculate beliefs at the very end.
  - *BeliefUpdate* has beliefs from the start, and keeps them around the whole time.

# Complexity

- Linear in number of cliques, total size of all factors.
- Can accomplish convergence in the same upward/downward passes we used for the earlier version.

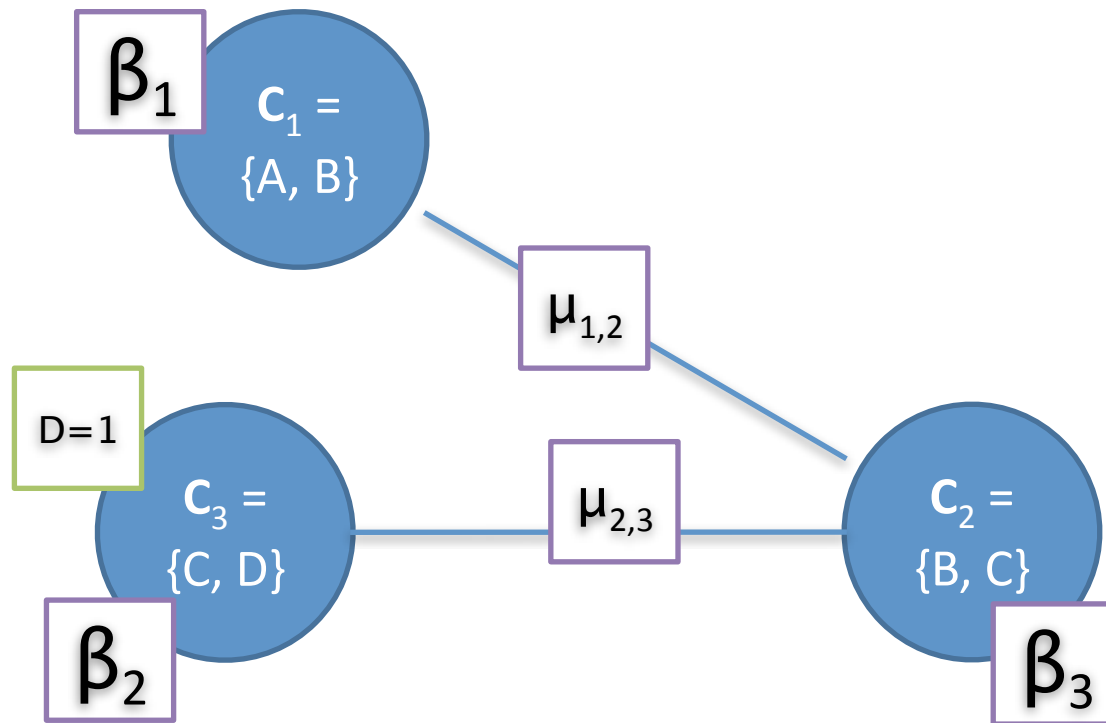
# Dealing with Evidence

# Advantage: Incremental Updates

- Naively, if we have evidence, we can alter the initial potentials at the start, then calibrate using message passing.
- Better: think of evidence as a newly arrived factor into some clique tree node(s).

# Example

$$\begin{aligned}\sigma_{i \rightarrow j} &\leftarrow \sum_{C_i \setminus S_{i,j}} \beta_i \\ \beta_j &\leftarrow \beta_j \times \frac{\sigma_{i \rightarrow j}}{\mu_{i,j}} \\ \mu_{i,j} &\leftarrow \sigma_{i \rightarrow j}\end{aligned}$$



# Advantage: Incremental Updates

- Naively, if we have evidence, we can alter the initial potentials at the start, then calibrate using message passing.
- Better: think of evidence as a newly arrived factor into some clique tree node  $i$ .
- Recalibrate: pass messages out from node  $i$ .  
Single pass!
- Retraction: can't recover anything multiplied by zero.



# Queries across cliques

# Advantage: Queries across Cliques


- Naively: enforce that all query variables are in some clique.
  - Every query might need its own clique tree!
- Better: variable elimination in a calibrated clique tree.
  - Bonus: only have to use a subtree that includes all query variables.

# Multi-Clique Queries

- Find a subtree of  $T$  that includes all query variables  $\mathbf{Q}$ . Call it  $T'$  and its scope  $\mathbf{S}$ .
- Pick a root node  $r$  in  $T'$ .
- Run variable elimination of  $\mathbf{S} \setminus \mathbf{Q}$  with factors (for all  $I$  in  $T'$ ):

$$\phi_i = \frac{\beta_i}{\mu_{i,\text{upstream}(i)}}$$

# Example: $Z\text{-}P(B, D)$



# Advantage: Multiple Queries


- Suppose we want the marginal for every *pair* of variables  $X, Y$ .
- Naïve: construct a clique tree so all nodes pair together. (Very bad.)
- Naïve: run VE  $n$ -choose-2 times.
- Better: dynamic programming.

# Dynamic Programming for All Pairs

- Construct a table so that  $A_{i,j}$  contains  $U(\mathbf{C}_i, \mathbf{C}_j) = Z \cdot P(\mathbf{C}_i, \mathbf{C}_j)$ .
- Base case:  $\mathbf{C}_i$  and  $\mathbf{C}_j$  are neighboring cliques.

$$\begin{aligned} A_{i,j} &= U(\mathbf{C}_i, \mathbf{C}_j) \\ &= U(\mathbf{C}_j \mid \mathbf{C}_i) U(\mathbf{C}_i) \\ &= \frac{\beta_j}{\mu_{i,j}} \beta_i \end{aligned}$$


- Proceed to farther more distant pairs recursively.



# Dynamic Programming for All Pairs

- $\mathbf{C}_i$  and  $\mathbf{C}_j$  are independent given  $\mathbf{C}_l$ .
- We already have  $U(\mathbf{C}_i, \mathbf{C}_l)$  and  $U(\mathbf{C}_l, \mathbf{C}_j)$ .

$$\begin{aligned} A_{i,j} &= U(\mathbf{C}_i, \mathbf{C}_j) \\ &= \sum_{\mathbf{C}_l \setminus \mathbf{C}_j} U(\mathbf{C}_i, \mathbf{C}_l) U(\mathbf{C}_j \mid \mathbf{C}_l) \\ &= \sum_{\mathbf{C}_l \setminus \mathbf{C}_j} A_{i,l} \frac{\beta_j}{\mu_{j,l}} \end{aligned}$$



# Pros and Cons:

## Message Passing in Clique Trees

- Multiple queries
- Incremental updates
- Calibration operation has transparent complexity.

But:

- Complexity can be high (space!)
- Slower than VE for a single query
- Local factor structure is lost



# Summary:

## Message Passing in Clique Trees

- How to construct a clique tree (from VE elimination order or triangulated chordal graph)
- Marginal queries for *all* variables solved in only twice the time of one query!
- Belief update version: clique potentials are reparameterized so that the clique tree invariant always holds.
- Runtime is linear in number of cliques, exponential in size of the largest clique (# variables; induced width).